



Artificial Intelligence II

Lesson 6- Learning Methods



First Code
Academy



Today's Plan

Teach Back	00 - 5 min
Supervised Learning	10 - 15 min
Reinforcement Learning	15 - 20 min
Quiz	20-25 min
Break	25 - 28 min
Project - Game Training	28 - 55 min
Lesson Recap	55 - 60 min



Teach Back

_____ allows us to easily find how variables relate to each other

The algorithm we use to improve our model is called gradient _____



What did we learn last time?



Regression

Given many (x, y) points,
find the relation between
the two variables


We need the **slope** and
intercept in this case





Gradient Descent

Gradient Descent finds the optimal parameters by moving in the **descending** direction





Key Terms

Supervised Learning

Reinforcement Learning



Supervised Learning



Supervised Learning

Supervised Learning means that we know the intended answer during training.

We have a **reference** answer during training.





Supervised Learning

Suppose we want to see if
a picture contains a cat.

The training data would
say if there is a cat or not.

The data is “**labeled**”

[cat]





Supervised Learning

In our regression project, the y in (x, y) was our label.

We saw how accurate our model was

Labeled Training Data

(x, y)

Our model's data

(x, \hat{y})

↑ The **error** is their difference ↓



Labeling Data

However, labeling data can be a time-consuming and expensive process

A lot of it is done by hand!



Amazon Mechanical Turk

Marketplace where companies can put up jobs such as labeling

Workers compete for these jobs.



Reinforcement Learning



Reinforcement Learning

Instead of using labeled data, we think in terms of **rewards** and **punishments**



Reinforcement Learning


Components in reinforcement learning:

1. **State:** Current environment conditions.
2. **Action:** A possible move.
3. **Reward:** The value in taking a certain action.



Reinforcement Learning

We want to take the actions with the highest reward



Which way should Pacman go?

Each **action** gives us a different value



Exploration vs. Exploitation

Our model wants to **explore** new strategies but also **exploit** the reward

An optimal strategy might involve short-term losses



Q-Learning

In Q-Learning, we have a function that gives us the “quality” of taking an action at a state


$$Q(s, a)$$

If at state s we take action a , what is the value of **all the following actions?**



Q-Learning

If we take **this** path, we would have more moves left in total, so higher value!





Quiz: bit.ly/FCA_AI_Quiz




Project : Game Training



Objective

We want to create an agent that learns to play a simple game using Q-Learning!



Our agent has to learn to balance a pole on a moving cart



OpenAI Gym

Gym is a Python library that allows us to learn reinforcement learning in games!

Install it on your machine by typing:

```
pip install gym
```




Get the starter file

http://bit.ly/FCA_AI2_Starter



**Make an agent that chooses
random action**



Random action

`render()` shows the current frame on the screen

`step()` takes an action from set of possible actions.


Always add `env.close()` at the very end!

```
1  import gym
2
3  # Create the environment
4  env = gym.make('CartPole-v0')
5
6  # Reset the environment
7  env.reset()
8  for _ in range(1000):
9      # show the current state on the screen
10     env.render()
11
12     # Take a random action
13     env.step(env.action_space.sample())
14 env.close()
```



Random action

Our agent doesn't react to its environment and just moves randomly!





Let's code a smart agent!



Creating agent

First we create the agent and environment

```
42     # Create the environment
43     env = gym.make('CartPole-v0')
44     # Create our agent
45     agent = CartPole(env)
```



Outer Loop

Set variables for every iteration of the game.

```
for ep in range(EPIISODES):  
    print(f'Episode: {ep}')  
  
    # Get the initial state  
    observation = env.reset()  
    # Turn continuous state space into discrete  
    state = agent.discretize(observation)  
  
    # Get Learning and Exploration rates  
    alpha = agent.get_alpha(ep)  
    epsilon = agent.get_epsilon(ep)  
  
    done = False  
    i = 0
```



Inner Loop

Handle the main game-playing logic

```
60         i = 0
61         while not done:
62             # Render the frame
63             env.render()
64
65             # Choose the action
66             action = agent.get_action(state, epsilon)
67
```




Inner Loop

Take the step and update our Q-table

```
68 # Take the action in game
69 observation, reward, done, _ = env.step(action)
70 new_state = agent.discretize(observation)
71
72 # update our Q-Table
73 agent.update_q(state, action, reward, new_state, alpha)
74 state = new_state
75 i += 1
```



CartPole Agent Class

We then setup the CartPole agent class

```
11 class CartPole():
12     def __init__(self, env, buckets=(1, 2, 6, 12)):
13         # Making the ranges into discrete ranges
14         self.buckets = buckets
15         self.env = env
16
17         # Q table
18         self.Q = np.zeros(self.buckets + (self.env.action_space.n,))
19
```



Get Action

Update the `get_action()` method.

```
35 def get_action(self, state, epsilon):  
36     return self.env.action_space.sample() if (np.random.random() <= epsilon) else np.argmax(self.Q[state])  
37
```



Update Q-Table


After each state transition, we update our table

```
38 def update_q(self, state, action, reward, new_state, alpha):
39     self.Q[state][action] += alpha * (reward + GAMMA * np.max(self.Q[new_state]) - self.Q[state][action])
40
```



Try it out!

Our agent learns slowly, so it will take a while for it to learn to play.





That's it for today!



Key Terms

Supervised Learning

Reinforcement Learning



Artificial Intelligence II

Lesson 6 - Learning Methods



First Code
Academy